

DATA VIRTUALITY MASTERCLASS

Topic: Kafka/KSQLDB showcase

Welcome to the DV Masterclass! Agenda



2:00pm - 3:00pm: Retrieving Data from Data Virtuality via ODATA

3:00pm - 3:15pm: Break

3:15pm - 4:15pm: Kafka/KSQLDB showcase

4:15pm - 4:20pm: Break

4:20pm - 5:15pm: Procedural Relational Command

What to expect from this session?

In this track, we will show how to get data from Kafka

- What is Kafka
- What is ksqlDB
- Use cases
- Options to connect
- KSQL SQL syntax specialties and Translator
- Queries from DV

What is Kafka?

What is Kafka?

- Kafkas 3 key abilities:
 - To publish (**write**) and subscribe to (**read**) **streams of events**, including continuous import/export of your data from other systems.
 - To store streams of events durably and reliably for as long as you want.
 - To process streams of events as they occur or retrospectively.
- Terminology: **events** are stored in **topics**, configurable how long
 - Topics are partitioned
- Challenge for Data Virtualization:
 - events are active, so we should react, but DV doesn't care if no one is "asking"
 - new events are coming in while a query is running
 - events will not be available after retention period

Use cases for Data Virtualization

Use cases for Data Virtuality

What would be your use cases? Please write them into the chat now!

Use cases from [Apache Kafka itself](#):

- **Messaging**
- **Website Activity Tracking**
- **Metrics (operational monitoring)**
- **Log Aggregation**
- **Stream Processing**
- **Event Sourcing**
- **Commit Log**

Community use cases:

Sebastian wrote: Continuous materialization of changes in DV

Use cases for Data Virtuality



Commented use cases

- **Messaging** - sending messages from a DV automation is possible, great use case
- **Website Activity Tracking** - classical use case, we would probably join this to our data, in an analytical way
- **Metrics** - Possible in a pull configuration, also emit metrics from DV
- **Log Aggregation** - Abstraction for Log files
- **Stream Processing** - IMHO very large use case, as we can integrate the results of processed and aggregated streams
- **Event Sourcing** - *is a style of application design where state changes are logged as a time-ordered sequence of records. Kafka's support for very large stored log data makes it an excellent backend for an application built in this style.*
- **Commit Log** - *Kafka can serve as a kind of external commit-log for a distributed system.*


What is ksqlDB

What is ksqlDB?

- ksqlDB is part of the confluent platform (Kafka distribution with extras)
- Provides (among other things) an SQL bridge to kafka, ideal for data virtualization
- ksqlDB is a stream processing database

capabilities:

- Derive a new stream from an existing stream
- Derive a new table from an existing stream
- Derive a new table from an existing table
- Derive a new stream from multiple streams



Streams and Tables

- A **stream** is an immutable, append-only collection that represents a series of historical facts, or events. Once a row is inserted into a stream, the row can never change. You can append new rows at the end of the stream, but you can't update or delete existing rows.
- A **table** is a mutable collection that models change over time. It uses row keys to display the most recent data for each key. All but the newest rows for each key are deleted periodically. Also, each row has a timestamp, so you can define a windowed table which enables controlling how to group records that have the same key for stateful operations – like aggregations and joins – into time spans. Windows are tracked by record key.

ksqlDB Query Types



- **Persistent query:** Persistent queries are server-side queries that run indefinitely processing rows of events.
- **Push query:** A push query is a form of query issued by a client that subscribes to a result as it changes in real-time.
- **Pull query:** A pull query is a form of query issued by a client that retrieves a result as of "now", like a query against a traditional RDBMS.

Options to connect

Options to connect from DV

- Via REST API
- Via generic JDBC with <https://github.com/mmolimar/ksql-jdbc-driver>
- Via DV driver implementing <https://docs.ksqldb.io/en/latest/developer-guide/ksqldb-clients/java-client/>

Via REST API

- A very powerful way to interact with ksqlDB
- Push queries provide the challenge of a streaming result set: *Execute a push query by sending an HTTP request to the ksqlDB REST API, and the API sends back a chunked response of indefinite length.*
- So we want pull queries

Structured data import wizard

Select JSON/XML source

Select source, prototype node and proceed to the next step

File source | Web service

Data source: ws_generic

End point: http://localhost:8088/ksql

Action: POST | Request content: application/json | Encoding: UTF-8

Request headers:

Request body: { "ksql": "show streams;", "streamProperties": {} }

Response format: JSON

Request rate limit: 1

Apply

Row configuration | Source data preview

☒ Fetch items by selected path in the document: "/root"

☐ Fetch all items in the document: "/root"

Select a prototype node using context menu or double click: Each root node below is a row in generated table:

root

- root
- _u0040_type
- statementText
- streams
- type
- name
- topic
- keyFormat

root

- _u0040_type
- statementText
- streams
- streams
- streams
- streams
- streams

< Back Next > Finish Cancel

Structured data import wizard

Create output columns

Define what attributes and nodes of the XML object will be projected into the destination table, using context menu in the tree below.

Prototype XML fragment

Manage output columns

Column...	Type	Default	Path	Precision	Scale	Date for...
idCol...	INTEGER					
type	STRING		type	4000		
name	STRING		name	4000		
topic	STRING		topic	4000		
keyFo...	STRING		keyFor...	4000		
value...	STRING		valueFo...	4000		

	idColumn	type	name	topic	keyFormat	valueFormat
1	1	STREAM	KSQLE_PROCESSING_LOG	default_ksql_processing_log	KAFKA	JSON
2	2	STREAM	PAGEVIEWS_REGION_LIKE_89	pageviews_filtered_r8_r9	KAFKA	AVRO
3	3	STREAM	USER_PAGEVIEWS	USER_PAGEVIEWS	KAFKA	AVRO
4	4	STREAM	PAGEVIEWS_STREAM	pageviews	KAFKA	AVRO
5	5	STREAM	RIDERLOCATIONS	locations	KAFKA	JSON

< Back Next > Finish Cancel

Via JDBC: Create a ksqlDB Connection

- Connector implements the java client in <https://docs.ksqldb.io/en/latest/developer-guide/ksqldb-clients/java-client/>
The client supports pull and push queries; inserting new rows of data into existing ksqlDB streams; creation and management of new streams, tables, and persistent queries; and also admin operations such as listing streams, tables, and topics.

```
call SYSADMIN.createConnection ( 'ksqldb', 'ksqldb', 'host=localhost,port=8088' );
call SYSADMIN.createDatasource ( 'ksqldb', 'ksqldb', '', 'supportsNativeQueries=true' );
```

- Metadata is being loaded.

Name	Description				
KSQLE_PROCESSING_LOG					
PAGEVIEWS_PER_REGION_89					
PAGEVIEWS_REGION_LIKE_89					
PAGEVIEWS_STREAM					
USER_PAGEVIEWS					
USERS_TABLE					

Columns					
Columns: 4					
Name	Type	Size	Scale	Nullable	Minimal
USERID	string	2147483647	0	NULL	
GENDER	string	2147483647	0	NULL	
REGIONID	string	2147483647	0	NULL	
NUMUSERS	bigint	2147483647	0	NULL	

Demo: Creating and Querying objects in ksqlDB

Create and query ksqlDB Stream

We are following the confluent platform quickstart

<https://docs.confluent.io/platform/current/platform-quickstart.html#quick-start-for-cp>

```
call "ksqldb.native"(  
  'CREATE or replace STREAM pageviews_stream WITH (KAFKA_TOPIC='pageviews',  
  VALUE_FORMAT='AVRO');', 'stmt');
```

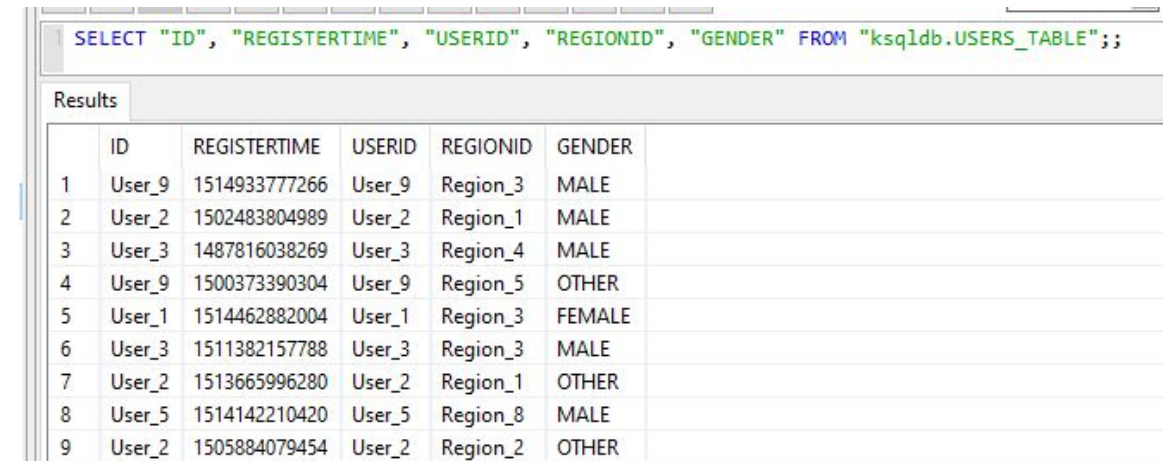
note the second parameter, it can either be **stmt**, **query** or **stream**

```
-- direct select  
SELECT "VIEWTIME", "USERID", "PAGEID" FROM "ksqldb.PAGEVIEWS_STREAM" LIMIT 10;;  
  
-- native way:  
select  
  *  
  from  
    ( call "ksqldb.native" ( 'SELECT VIEWTIME, USERID, PAGEID FROM pageviews_stream EMIT  
CHANGES LIMIT 10;', 'query' ) ) a  
    , ARRAYTABLE ( a.tuple COLUMNS "VIEWTIME" bigint, "USERID" string, "PAGEID" string ) as w;;  
  
select  
  cast (tuple as string)  
  from  
    ( call "ksqldb.native" ( 'SELECT VIEWTIME, USERID, PAGEID FROM pageviews_stream EMIT  
CHANGES LIMIT 10;', 'query' ) ) a
```

Create and query ksqlDB Table

```
call "ksqldb.native"(  
  'CREATE OR REPLACE TABLE users_table (id VARCHAR PRIMARY KEY)  
  WITH (KAFKA_TOPIC='users', VALUE_FORMAT='AVRO');', 'stmt');
```

- Tables in ksqlDB need a primary key!
- Result will be the same for each query
- Result will be in fetching mode in DV until cancelled



The screenshot shows a query execution interface. At the top, a SQL query is entered: `SELECT "ID", "REGISTERTIME", "USERID", "REGIONID", "GENDER" FROM "ksqldb.USERS_TABLE";`. Below the query, a tab labeled "Results" is active. The results are displayed in a table with 6 columns: ID, REGISTERTIME, USERID, REGIONID, GENDER, and an empty column. There are 9 rows of data.

	ID	REGISTERTIME	USERID	REGIONID	GENDER	
1	User_9	1514933777266	User_9	Region_3	MALE	
2	User_2	1502483804989	User_2	Region_1	MALE	
3	User_3	1487816038269	User_3	Region_4	MALE	
4	User_9	1500373390304	User_9	Region_5	OTHER	
5	User_1	1514462882004	User_1	Region_3	FEMALE	
6	User_3	1511382157788	User_3	Region_3	MALE	
7	User_2	1513665996280	User_2	Region_1	OTHER	
8	User_5	1514142210420	User_5	Region_8	MALE	
9	User_2	1505884079454	User_2	Region_2	OTHER	

Pull Queries

- Using the KSQLDB directly, see below challenge
 - `select * from USERS_TABLE WHERE REGISTERTIME > 1488520368669;`
 - *WHERE clause missing key column for disjunct: (REGISTERTIME > 1488520368669). See <https://cnfl.io/queries> for more info. Add EMIT CHANGES if you intended to issue a push query. Pull queries require a WHERE clause that: - includes a key equality expression, e.g. `SELECT * FROM X WHERE <key-column>=Y`; - in the case of a multi-column key, is a conjunction of equality expressions that cover all key columns. If more flexible queries are needed, table scans can be enabled by setting `ksql.query.pull.table.scan.enabled=true`. Statement: `select * from USERS_TABLE WHERE REGISTERTIME > 1488520368669;`*
 - `select * from USERS_TABLE where ID='User_3';`
 - *The `USERS_TABLE` table isn't queryable. To derive a queryable table, you can do 'CREATE TABLE QUERYABLE_USERS_TABLE AS SELECT * FROM USERS_TABLE'. See <https://cnfl.io/queries> for more info. Add EMIT CHANGES if you intended to issue a push query. Statement: `select * from USERS_TABLE Where ID='User_3';`*

But our connector takes care of this:

```
SELECT "ID", "REGISTERTIME", "USERID", "REGIONID", "GENDER" FROM "ksqldb.USERS_TABLE" WHERE REGISTERTIME > 1488520368669
```

Results						
	ID	REGISTERTIME	USERID	REGIONID	GENDER	
1	User_9	1514933777266	User_9	Region_3	MALE	
2	User_2	1502483804989	User_2	Region_1	MALE	
3	User_9	1500373390304	User_9	Region_5	OTHER	
4	User_1	1514462882004	User_1	Region_3	FEMALE	
5	User_3	1511382157788	User_3	Region_3	MALE	
6	User_3	1513666666666	User_3	Region_3	OTHER	

Create and query a joined stream

```
call "ksqldb.native"(  
  'CREATE STREAM user_pageviews  
  AS SELECT users_table.id AS userid, pageid, regionid, gender  
    FROM pageviews_stream  
    LEFT JOIN users_table ON pageviews_stream.userid = users_table.id EMIT CHANGES;', 'stmt');;  
  
;;  
  
SELECT "USERID", "PAGEID", "REGIONID", "GENDER" FROM "ksqldb.USER_PAGEVIEWS" limit 10;;
```

- JOIN operation is happening on the ksqlDB side

Further operations: Filtered Stream

```
call "ksqldb.native"(  
  'CREATE OR REPLACE STREAM pageviews_region_like_89  
  WITH (KAFKA_TOPIC='pageviews_filtered_r8_r9', VALUE_FORMAT='AVRO')  
  AS SELECT * FROM user_pageviews  
  WHERE regionid LIKE '%_8' OR regionid LIKE '%_9'  
  EMIT CHANGES;', 'stmt');
```

- Derived stream
- ksqlDB processing

Further operations: Windowed Table

```
call "ksqldb.native"(  
  'CREATE TABLE pageviews_per_region_89 WITH (KEY_FORMAT='JSON')  
  AS SELECT userid, gender, regionid, COUNT(*) AS numusers  
    FROM pageviews_region_like_89  
    WINDOW TUMBLING (SIZE 30 SECOND)  
    GROUP BY userid, gender, regionid  
    HAVING COUNT(*) > 1  
  EMIT CHANGES;', 'stmt');
```

- Here, we look into the last 30 seconds of events
- Persistent

Creating Tables / Writing into Streams

```
call ks.native('create table test1(id integer primary key, a varchar) WITH  
(kafka_topic='locations', value_format='json', partitions=1);','stmt');
```

```
call ks.native('insert into test1 values(1, 'aaa');','stmt');
```

```
call "ksqldb.native"('INSERT INTO riderLocations (profileId, latitude, longitude) VALUES  
(\'c2309eec\', 37.7877, -122.4205);','stmt');
```

- This enables you to write data from DV into a kafka stream

Any feedback / questions?

Thank you!

Please feel free to contact us at:
presales@datavirtuality.com

or

visit us at:
datavirtuality.com